

From Legacy Infrastructure to Sovereign Enforcement: A Kubernetes-Rust Blueprint for Safe Neuromorphic Evolution

The Kubernetes Sovereign Control Plane: Designing a Neuro-Safe Orchestration Substrate

The strategic objective is to repurpose Kubernetes, Helm, and Prometheus not merely as infrastructure but as a sovereign enforcement layer for next-generation neuromorphic systems. This paradigm shift leverages Kubernetes' mature abstractions for policy enforcement, state management, and declarative configuration to create tangible, non-negotiable guardrails for computational workloads operating within biological and ethical boundaries [20](#) [24](#). By treating Kubernetes as a foundational cybernetic laboratory, it becomes the central nervous system for managing evolution proposals, enforcing neurorights, and monitoring real-time safety metrics like the Rate of Harm (RoH). This approach provides real enforcement hooks that reject any deployment violating formal safety rules before it impacts the cluster, thereby establishing a verifiable audit trail of all system changes [18](#) [19](#). The sovereignty layer is realized through a set of custom resources (CRDs) that serve as both the schema for configuration and the unit of enforcement, complemented by powerful admission controllers that act as the system's conscience.

The cornerstone of this sovereign layer is the creation of a suite of Kubernetes CRDs designed to mirror the internal data structures of the neuromorphic stack, such as `.ocpu`, `.ocpuenv`, and `.aln` files. These CRDs translate abstract safety concepts into first-class, versionable objects within the Kubernetes API. The primary CRDs for this system are `OrganicCpuProfile`, `OrganicCpuEnvelope`, `NeuroRightsPolicy`, `EvolutionProposal`, and `EvidenceBundle`.

CRD Name	Purpose	Key Fields
OrganicCpuProfile	Defines the static capabilities and personality of a neuromorphic worker. It serves as a blueprint for safe computational modes.	<code>bio_limits</code> : Compute envelope (G, D, etc.). <code>aln_grammar</code> : Embedded ALN particles defining valid actions. <code>neurorights_policy_ref</code> : Reference to a <code>NeuroRightsPolicy</code> object.
OrganicCpuEnvelope	Represents the dynamic, real-time constraints on a running workload. It reflects the current state of the system.	<code>lifeforce</code> : Cy/zen/chi levels. <code>roh</code> : Rate of Harm. <code>cognitive_load</code> : Current duty cycle. <code>violations</code> : Status indicating breached limits.
NeuroRightsPolicy	A declarative specification of ethical constraints, potentially scoped to different jurisdictions to enforce a "strictest-wins" rule.	<code>jurisdiction</code> : K8s Namespace scope. <code>rights_clauses</code> : Neuroright definitions (e.g., freedom of thought, mental privacy). <code>duty_cycles</code> : Enforced behavioral patterns.
EvolutionProposal	The central artifact for proposing a change to the system. It represents a high-level intent to evolve the system safely.	<code>update_proposal</code> : Typed struct detailing changes (L2 delta, irreversibility flag). <code>evidence_bundle_ref</code> : Reference to an <code>EvidenceBundle</code> containing supporting data.
EvidenceBundle	A collection of data points used to justify an <code>EvolutionProposal</code> , enabling long-horizon learning and accountability.	<code>pre_post_envelopes</code> : Parameters before and after a proposed update. <code>roh_estimates</code> : Estimated Rate of Harm values. <code>simulation_logs</code> : Links to replayable logs (.ocpulog).

The true power of this design lies in the admission controller logic that governs these resources. An admission controller is a piece of code that intercepts requests to the Kubernetes API server prior to persistence of the resource, allowing for validation and mutation [24](#). In this context, a validating admission webhook would be configured to inspect every `EvolutionProposal` and `Deployment` using an `OrganicCpuProfile` [64](#) [67](#). This webhook would perform a multi-stage validation process. First, it conducts a standard schema validation, ensuring the incoming YAML or JSON conforms to the OpenAPI v3 schema for the respective CRD, preventing malformed inputs [39](#) [56](#). Second, it performs a static safety check, comparing the proposed `UpdateProposal` against the immutable `BioLimits` in the associated `OrganicCpuProfile`. Any proposal attempting to violate monotonic invariants, such as decreasing capability (`G_new < G_old`) or increasing degradation (`D_new > D_old`), would be immediately rejected .

Third, the admission controller executes a critical dynamic safety check by querying Prometheus for the current state of the target workload(s). Prometheus acts as a bioscale monitor, exporting metrics derived from the Rust crates, including RoH, knowledge-factor, fatigue indices, and lifeforce envelopes . The admission controller would retrieve the corresponding `OrganicCpuEnvelope` and check if applying the proposed update would push any metric, such as RoH or lifeforce, beyond its predefined thresholds. For instance, if the RoH alert rule is set to "`RoH > 0.3`", the admission controller would query Prometheus for the current RoH value of the target workload; if it is already at 0.29 and the proposed update is expected to raise it further, the proposal is blocked . This mechanism directly implements the fail-safe design principles advocated by standards

like IEEE P7009 for autonomous systems, which emphasize establishing robust, transparent safeguards against misuse and manipulation [15](#) [16](#). The final and most complex validation step involves evaluating the proposal against ALN grammars and neurorights policies. To achieve this, the admission webhook would invoke a trusted service—likely a gRPC endpoint exposed by the `sovereigntycore` Rust crate—which evaluates the `UpdateProposal` against the active `NeuroRightsPolicy`. The proposal is only admitted if this external evaluation returns a `DecisionOutcome::Allowed`.

For practical implementation, while traditional admission webhooks are often written in Go, the project's focus on Rust suggests leveraging libraries like `kube-rs` to build the admission controller entirely in Rust [18](#) [19](#). This ensures type safety and performance consistency across the entire stack. Furthermore, to ensure the absolute correctness of this critical admission logic, formal methods can be applied. Frameworks exist that model Kubernetes authorization and admission policies as first-order logic and use SMT solvers like Microsoft's Z3 to exhaustively search for misconfigurations and logical flaws [20](#) [23](#). Research has demonstrated the use of Z3 for modeling and verifying frameworks like DeepGlobal and NRSMs, proving the feasibility of this approach [63](#) [132](#). Tools like Anvil, which uses the Verus verification framework, are being developed specifically to verify the correctness of Kubernetes controllers, reinforcing the viability of building a formally verified sovereignty layer [22](#) [37](#). This combination of declarative CRDs, imperative admission controllers, and formally verifiable logic creates a robust, auditable, and ultimately sovereign control plane.

The Rust Foundation: Embedding Compile-Time Safety Guarantees

While Kubernetes provides the runtime enforcement layer, the ultimate strength of the sovereign system resides in embedding these safety properties directly into the software fabric via the Rust programming language. The proposed set of Rust crates forms a cohesive, type-safe foundation that makes violations impossible at the application level. This approach moves safety from a runtime check to a compile-time guarantee, a principle central to Rust's philosophy. The key crates—`sovereigntycore`, `organiccpucore`, `organiccpualn`, `cybernano-viability-kernel`, and `cybernano-vector-cyberrank`—are designed to work in concert, creating a closed loop where declarations in Kubernetes are enforced by the execution logic in Rust.

The `sovereigntycore` crate is the central authority for all safety decisions. It exposes a public API, likely through a gRPC service, that other components can query. Its primary function would be `evaluate_update`, which takes an `UpdateProposal` and a `NeuroRightsPolicy` as input and returns a `DecisionOutcome` (e.g., `Allowed` or `Denied`). This crate is responsible for orchestrating the evaluation process, calling upon other specialized crates to perform specific checks. For example, it would delegate ALN grammar validation to the `organiccpcualn` crate and execute the viability kernel calculations using the `cybernano-viability-kernel` crate. The logic within `sovereigntycore` itself could be subjected to rigorous formal verification using tools like Verus, which allows developers to write mathematical proofs about their code's behavior, ensuring that the decision-making process is logically sound under all possible conditions ^{[26](#)}.

The `organiccpcucore` crate is the engine of the OrganicCPU. It defines the fundamental state and behavior of a neuromorphic worker. Crucially, its public API is designed around the concept of envelopes. Instead of exposing raw computational functions that could be misused, it provides operations that explicitly take an `OrganicCpuEnvelope` as a parameter. For example, a function to increase computational intensity might have a signature like `fn intensify(current_env: &Envelope) -> Result<Envelope, ViolationError>`. This function would calculate the new state based on the desired intensity and the current `Envelope`. If the resulting state violates any bounds (e.g., `RoH` exceeds 0.3 or `lifeforce` drops below a threshold), the function would return an error, making the violation a compile-time checked result or a controlled runtime failure, rather than an undefined state. This design enforces the principle that all workloads must operate within thermodynamically and neurorights-constrained envelopes, a core requirement of the hexen-unified-rs architecture.

The `organiccpcualn` crate's sole purpose is to parse and represent the `.aln` (Artificial Life Nucleotides) files, which contain the grammatical rules governing valid state transitions. By mapping these text-based files to strongly-typed Rust structs using a library like `serde`, this crate ensures that ALN grammars are syntactically and structurally sound before they are ever passed to the `sovereigntycore` for evaluation. This prevents a large class of errors related to malformed policy files. The parsed structures become the direct input for the viability kernel calculations, ensuring that the system's operational boundaries are always grounded in a formally defined grammar.

The `cybernano-viability-kernel` and `cybernano-vector-cyberrank` crates implement the advanced control algebra necessary for Tsafe decision-making. The `viability_kernel` component computes the "safe set"—the region in the 7-D state space (intensity, duty cycle, cumulative load, etc.) where the system can operate without

violating its constraints. This is visualized as a polytope in the Unreal Engine HUD . The CyberRank component then operates on this viable set, using Pareto optimization to select the best course of action from a candidate set of proposals. It balances competing objectives, such as maximizing capability gain (G) while minimizing risk factors like RoH and psych risk, presenting a ranked list of options to the operator . These two crates are the workhorses inside the `sovereigntycore`'s evaluation function, providing the sophisticated mathematical reasoning needed to make nuanced safety decisions.

Finally, the `hexen-backend-core` and `cyberswarm-neurostack` crates are built upon this foundation. They are wired to use the shared safety crates, ensuring that all neuromorphic code shares the same safety algebra . The `hexen-backend-core` crate acts as the entry point, exposing guard endpoints that receive high-level intents from AI-chats. These intents are translated into typed `UpdateProposals` and passed to the `sovereigntycore` for evaluation before any underlying action is taken . The `bioscale-upgrade-store` and `evidence-registry` crates become the persistent storage for the "donutloop," storing every evolution step along with its accompanying `EvidenceBundle`, which contains pre/post envelope parameters and RoH estimates for long-horizon learning and auditing . This tight coupling between the declarative Kubernetes layer and the imperative Rust execution layer creates a coherent, robust, and highly secure system where safety is a first-class citizen woven into every part of the architecture.

Structuring AI-Driven Evolution Proposals Through Machine-Readable Schema

The vision for AI-driven evolution requires a fundamental departure from free-form natural language prompts. To ensure safety and verifiability, AI-chat agents must produce structured, schema-compliant edits to the Kubernetes CRDs and underlying Rust crate configurations. This transforms the AI from a simple prompt generator into a sophisticated editor for a formal system. The key to achieving this is to provide the AI with access to the complete, machine-readable schemas of the system's core artifacts, allowing it to auto-complete, validate, and generate correct YAML manifests instead of unstructured text. This approach aligns with modern agent communication protocols designed to bridge the gap between human intent and machine-executable commands

The primary mechanism for enabling this structured interaction is the Model Context Protocol (MCP), an emerging standard for LLM-powered applications to interact with complex APIs [54](#) [62](#). An MCP server for Kubernetes, such as the one mentioned in the provided context, can expose the full OpenAPI schema of the custom resources to the AI agent [55](#). This allows the AI to explore the available resources (`OrganicCpuProfile`, `EvolutionProposal`, etc.), understand their fields, types, and relationships, and generate compliant JSON or YAML representations without being overwhelmed by context length limitations [36](#). The AI agent would be prompted with a high-level intent, such as "increase difficulty of the neuromorphic scenario by 10%." The agent's task is then to translate this intent into a concrete `EvolutionProposal` CRD manifest. This manifest would be a well-formed YAML file specifying the desired changes, referencing a new `OrganicCpuProfile` that defines the increased difficulty, and linking to an `EvidenceBundle` containing the supporting data.

To facilitate this process, a small set of canonical, text-based formats with explicit schemas must be established for all key artifacts. These include `.ocpu` and `.ocpuenv` for CPU profiles and envelopes, `.aln` for ALN grammars, and `.evolve.jsonl` for evolution manifests. Each format would have a corresponding JSON Schema or CDDL definition. For example, the schema for an `EvolutionProposal` would define the structure of the `UpdateProposal` object, including fields for effect bounds (like an L2 delta) and an irreversible flag, and would reference the `EvidenceBundle` object [36](#). Libraries in Rust, such as `psx-api`, exist to validate responses against OpenAPI schemas, demonstrating the feasibility of building tooling to check the AI's output against these defined schemas [39](#). This validation step would occur in the CI/CD pipeline or even as part of the AI interaction flow itself, rejecting any generated manifest that fails to conform to the schema.

The workflow for an AI-driven evolution proposal would proceed as follows:

- 1. Intent Capture:** An operator provides a high-level intent to an AI chat interface.
- 2. Schema-Aware Generation:** The AI agent, equipped with the MCP server's schema information, drafts a YAML manifest for a new `EvolutionProposal` CRD. It auto-completes field names and validates the types of values it inserts.
- 3. Evidence Bundle Creation:** As part of the proposal, the AI would also generate or reference an `EvidenceBundle`. This bundle could include links to simulation results, updated `.ocpuenv` files showing expected state changes, and a justification for why the change adheres to neurorights constraints, formatted as a `.evolve.jsonl` file.
- 4. Human-in-the-Loop Review:** The generated YAML manifest and evidence bundle are presented to a human operator for review. The operator can inspect the proposed changes, the evidence, and the potential impact on safety metrics before authorizing submission.
- 5. Submission and Admission:**

Upon approval, the human operator applies the YAML manifest to the Kubernetes cluster. The admission controller then intercepts the request, performing its multi-faceted validation against the `OrganicCpuProfile`, current `OrganicCpuEnvelope` from Prometheus, and the `NeuroRightsPolicy` ²⁴. 6. **Execution or Rejection:** If the proposal passes all checks, the Kubernetes controller associated with the `EvolutionProposal` CRD begins executing the change, perhaps by deploying a new version of the `hexen-backend-core` with the updated configuration. If it fails any check, the proposal is rejected, and the reason for rejection is logged, providing valuable feedback for future attempts.

This structured workflow ensures that all evolution proposals are explicit, traceable, and subject to automated, rigorous safety checks. It leverages the strengths of AI for drafting and pattern recognition while retaining essential human oversight, creating a collaborative and safe environment for advancing neuromorphic systems. This approach is heavily inspired by NeuroPC/NeuroAutomatic patterns, which aim to reduce cognitive load by automating the translation of high-level intentions into low-level, safe commands .

Code Quality and Formal Verification: From Property Testing to Proof-Carrying Code

Improving code quality for neuromorphic stacks requires a symbiotic relationship between standardized file formats and formal safety properties. The goal is to move beyond syntactic correctness and conventional testing toward mathematical proof of safety, where builds and deployments fail if core safety axioms are violated. This is achieved through a layered defense-in-depth strategy encompassing standardized schemas, automated property-based testing in the CI/CD pipeline, and, for the highest assurance, formal verification of critical components.

The first layer is the establishment of standardized, text-based formats with explicit schemas for all configuration artifacts. This includes `.ocpu` (Organic CPU Profile), `.ocpuenv` (Envelope), `.ocpulog` (Log), `.aln` (ALN Grammar), and `.evolve.jsonl` (Evolution Manifest) . Defining these schemas is paramount because it enables machine-assisted validation at every stage of the development lifecycle. For instance, a JSON schema for an `OrganicCpuEnvelope` would precisely define the allowed keys (`lifeforce`, `roh`, `g_factor`, `d_factor`) and their data types, ensuring that no malformed or semantically incorrect data is introduced into the system. Validation tools

exist in many ecosystems, including Rust libraries like `psx-api` that can check response bodies against OpenAPI schemas, and command-line tools that can validate JSON files against a given schema [39](#) [45](#). By integrating schema validation into the pre-commit hooks and CI/CD pipelines, a vast class of bugs related to misconfiguration can be caught immediately, before they ever reach the Kubernetes API [11](#).

The second layer is a robust CI/CD pipeline that incorporates property-based testing. Unlike unit tests that check specific examples, property-based tests generate a wide range of random inputs to verify that a function holds true under all conditions. Using a library like `proptest` in Rust, one can write tests that enforce the system's formal safety properties [44](#). For example, a property test can be written to verify the monotonicity of OTA updates. Given a valid `EvolutionProposal` that has passed the admission controller's checks, the test would assert that the resulting `OrganicCpuEnvelope` satisfies the invariants $G_{new} \geq G_{old}$ and $D_{new} \leq D_{old}$. Another property test could verify that any accepted update never causes the calculated Rate of Harm (RoH) to exceed its specified bound (e.g., 0.3). These tests run automatically on every commit, providing a continuous check on the system's logical integrity. If a change breaks one of these hard-coded safety properties, the CI pipeline fails, blocking the unsafe code from being merged or deployed. This practice is a form of lightweight formal verification, using computational exploration to find counterexamples to safety claims.

The third and most rigorous layer is formal verification, which aims to mathematically prove that a piece of code meets its specifications. This is particularly suited for the `sovereigntycore` crate and the admission controller logic, as their correctness is paramount to the entire system's security. Languages and tools like Verus for Rust allow developers to write formal specifications (preconditions, postconditions, and invariants) alongside their code [26](#). The Verus compiler then generates proof obligations that it attempts to discharge using an SMT solver like Z3 [25](#) [26](#). If the proofs succeed, the code is guaranteed to meet its specifications for all possible inputs. Research has successfully applied this technique to verify the correctness of Kubernetes controllers and various neural network models, demonstrating its applicability to this domain [22](#) [37](#) [83](#). While more labor-intensive, formal verification provides the strongest possible guarantee of safety, eliminating entire classes of logical errors that might evade even exhaustive property testing.

The synergy between these layers creates a comprehensive quality assurance framework. Standardized schemas provide the vocabulary and syntax. Property-based testing provides a probabilistic check of the semantic meaning of that vocabulary. And formal verification provides a definitive, mathematical proof of the most critical parts of the

system's logic. This integrated approach tightly couples code quality with formal safety, ensuring that the resulting system is not just functional but provably safe according to the defined neurorights and bioscale constraints.

Layer	Technique	Tooling / Methodology	Objective
Layer 1: Schema Validation	Static validation against predefined schemas.	JSON Schema, OpenAPI v3, Command-line validators.	Prevent malformed inputs and ensure structural correctness of <code>.ocpu</code> , <code>.evolve.jsonl</code> , etc. 39 56
Layer 2: Property-Based Testing	Automated generation of random test cases to falsify safety properties.	Rust <code>proptest</code> library, QuickCheck.	Prove that invariants ($G_{new} \geq G_{old}$, $\text{RoH} \leq 0.3$) hold for all valid inputs. 44
Layer 3: Formal Verification	Mathematical proof of correctness using SMT solvers.	Verus for Rust, Anvil, Z3.	Provide a definitive guarantee that critical components (<code>sovereigntycore</code>) meet their safety specifications. 22 26 37

This multi-layered strategy ensures that the path from developer intent to deployed system is paved with verifiable safety, transforming the CI/CD pipeline from a simple build-and-deploy mechanism into a powerful engine for safety assurance.

System Integration and Read-Only Visualization with Unreal Engine

The architectural design positions Unreal Engine as a secondary, read-only visualization layer, acting as a safety-aware Heads-Up Display (HUD) for operators interacting with the sovereign neuromorphic system. Its role is not to control or modify the system's state but to provide intuitive, immersive visualizations of the complex safety metrics and control algebras computed by the Rust backend. This transforms abstract numerical data and theoretical concepts like viability kernels into an understandable spatial representation, enhancing operator awareness and trust without introducing any enforcement logic. All communication from Unreal to the Rust backend must be read-only, for instance, via FFI, HTTP, or gRPC endpoints.

The integration begins on the Rust side, within the `hexen-unified-rs` or `OrganicCPU` components. These services must compute and expose a rich set of state variables that encapsulate the system's health and operational boundaries. Based on the conversation history, this includes a 7-D SwarmState vector comprising metrics like intensity, duty cycle, cumulative load, neuromod amplitude, cognitive load, and legal complexity, as well as the three-part lifeforce metric (cy/zen/chi). These values are

continuously calculated and updated, reflecting the real-time status of the neuromorphic swarm. Alongside this state data, the Rust services must also compute the outputs of the advanced control algorithms, namely the viability kernel (a safe operating polytope in the 7-D state space) and the CyberRank vectors that rank candidate actions based on safety and efficacy .

These data points are exposed to Unreal Engine through well-defined, read-only interfaces. This could be implemented via a gRPC service exposed by the Rust backend, which Unreal plugins can call, or through a simpler HTTP/REST API. The endpoints would provide the current state, the vertices of the viability kernel polytope, the set of candidate actions being considered, their CyberRank scores, and the reasons for any actions being clipped or rejected (e.g., safety, rollback, psych risk) .

On the Unreal Engine side, custom C++ plugins and Blueprints would be developed to consume this data. The core components would be:

1. **UNeuropcViabilityKernelComponent:** This component would be attached to a scene element (e.g., a sphere or mesh representing the system). It would periodically fetch the 7-D viability kernel polytope from the Rust backend. Since a 7-D shape cannot be visualized directly, the component would project it onto a 3D subspace for rendering. This projection would give the operator a dynamic, intuitive sense of the "safe zone" within which the system must operate. The visualization could change color or shape as the system approaches the boundary of this kernel, providing a powerful visual warning.
2. **ANeuropcCyberRankController:** This actor would manage the visualization of candidate actions. When the `sovereigntycore` proposes several potential next steps, this controller would receive the corresponding CyberRank vectors. It would then render a Pareto front diagram, visually showing the trade-offs between different objectives (e.g., capability gain vs. harm reduction). The UI would clearly indicate why certain actions were accepted, clipped due to safety concerns, or rejected for other reasons like high psych risk, making the decision-making process transparent to the operator .

This Unreal-based HUD serves a crucial purpose in the context of "more-than-gaming" applications, particularly for age-gated content intended for users 25+ . It elevates the experience from passive consumption to active participation in a safe, sovereign microspace. The operator is not just playing a game; they are navigating a complex cybernetic system, guided by a visualization of its underlying safety logic. This direct link

to the control algebra and viability kernels, as described in the CyberNano spec, grounds the experience in a scientifically-inspired framework of safety and autonomy .

It is critical to maintain the read-only constraint strictly. Unreal Engine should have absolutely no permission to modify kernels, lifeforce envelopes, or neurorights policies. Any attempt by an external entity, including an Unreal plugin, to send modification requests to the Rust backend would be treated as a security violation and rejected. This ensures that the enforcement logic remains solely within the Kubernetes and Rust layers, while Unreal serves only as a sophisticated and engaging window into that secure system. This separation of concerns is vital for maintaining the integrity and sovereignty of the platform.

Synthesis and Research Roadmap

This research report has detailed a comprehensive architectural blueprint for a sovereign enforcement layer for neuromorphic computing. The core innovation lies in repurposing Kubernetes, Helm, and Prometheus from legacy infrastructure into a sophisticated cybernetic laboratory. This is achieved through a tightly integrated system of Kubernetes CRDs, a Rust foundation with compile-time safety guarantees, AI-driven evolution proposals structured by machine-readable schemas, and a rigorous, multi-layered code quality and verification process. The Unreal Engine serves as a secondary, read-only visualization layer, grounding the system in an intuitive user experience without compromising its security model.

The proposed system establishes a closed-loop governance cycle. Operators or AI agents propose an evolution via a structured `EvolutionProposal` CRD. The admission controller enforces static and dynamic safety rules, querying Prometheus for real-time metrics like RoH and lifeforce envelopes. A trusted `sovereigntycore` service, implemented in Rust, performs the final evaluation against ALN grammars and neurorights policies. Only if the proposal is deemed safe is it admitted into the cluster. The `hexen-unified-rs` and `OrganicCPU` backends, built upon the `organicccpu` and `cybernano-viability-kernel` crates, execute the change within the strict confines of the `OrganicCpuEnvelope`. Throughout this process, Prometheus monitors the outcome, and the entire event is recorded in the Kubernetes API, creating a verifiable audit trail. Finally, the Unreal Engine HUD provides a read-only visualization of the system's state and viability kernel, enhancing operator situational awareness.

Based on this analysis, the following research roadmap is recommended to bring this vision to fruition:

1. **Prioritize Core CRD and Admission Controller Design:** The immediate priority is to define the OpenAPI v3 schemas for the `OrganicCpuProfile`, `OrganicCpuEnvelope`, `NeuroRightsPolicy`, and `EvolutionProposal` CRDs in detail. Concurrently, develop a prototype admission controller using the `kube-rs` library in Rust. This prototype should focus on implementing the basic validation logic against the schemas and the monotonicity invariants ($G_{new} \geq G_{old}$, $D_{new} \leq D_{old}$).
2. **Define and Quantify Key Metrics:** A significant area of research is the formal quantification of the Rate of Harm (RoH) and the translation of neurorights into measurable constraints. This requires interdisciplinary work combining neuroscience, ethics, and computer science, drawing inspiration from frameworks like UNESCO's Recommendation on the Ethics of Neurotechnology [60](#) [65](#) [68](#). Developing a mathematical model for RoH is a prerequisite for writing meaningful property tests and configuring Prometheus alerts.
3. **Develop the `sovereigntycore` and `organiccpucore` Crates:** Begin implementing the foundational Rust crates. The `organiccpucore` crate should establish the `BioState`, `BioLimits`, and `OrganicCpuTick` concepts, with APIs that inherently respect the `OrganicCpuEnvelope`. The `sovereigntycore` crate should define the `UpdateProposal`, `NeuroRightsPolicy`, and `DecisionOutcome` structures and implement the initial logic for evaluating proposals against these types.
4. **Implement the CI/CD Verification Pipeline:** Integrate the nascent Rust crates into a CI/CD pipeline (e.g., using GitHub Actions [52](#)). Start with schema validation for the CRD manifests. Progress to implementing property-based tests using `proptest` to verify the safety invariants of the `organiccpucore` crate. This will provide early, automated feedback on the correctness of the safety logic.
5. **Explore Formal Verification for Critical Components:** Select the `sovereigntycore` crate for a pilot formal verification effort using a tool like Verus. The goal is to write formal specifications for its `evaluate_update` function and attempt to prove its correctness. This is a long-term goal that will provide the highest level of assurance but requires significant expertise.

6. Build the Unreal Engine Visualization Prototype: Develop a minimal Unreal Engine plugin that connects to the gRPC endpoints of the Rust backend. The initial prototype should focus on visualizing the 7-D SwarmState as a simple 3D scatter plot and displaying the current values of RoH and lifeforce metrics. This will validate the communication channel and provide insights into effective visualization techniques.

By systematically addressing these research items, the project can de-risk its ambitious goals and progressively build a truly sovereign and safe neuromorphic computing platform, where every line of code and every deployment is held accountable to a rigorous, verifiable standard of safety and ethics.

Reference

1. (PDF) Transistor → CPU → GPU → Model → Agent → ... https://www.researchgate.net/publication/396270471_Transistor_CPU_GPU_Model_Agent_Society_Autonomic_Systems_Objective_Layer_Co-Designed_Matter_Embodied_Mesh
2. A New Approach to Network Function Virtualization <https://cs.nyu.edu/~apanda/assets/dissertation/dissertation-updated.pdf>
3. <https://packages.debian.org/bookworm/ppc64/allpack...> <https://packages.debian.org/bookworm/ppc64/allpackages?format=txt.gz>
4. Neuromorphic Computing: A Theoretical Framework for ... https://www.researchgate.net/publication/393982506_Neuromorphic_Computing_A_Theoretical_Framework_for_Time_Space_and_Energy_Scaling
5. Neuromorphic Hardware Learns to Learn - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC6536858/>
6. What is Neuromorphic Computing? | Definition from ... <https://www.techtarget.com/searchenterpriseai/definition/neuromorphic-computing>
7. Neuromorphic Computing Gets Ready For the (Really) Big ... <https://cacm.acm.org/news/neuromorphic-computing-gets-ready-for-the-really-big-time/>
8. Advancing Neuromorphic Computing With Loihi: A Survey ... <https://www.researchgate.net/publication/>

350705246_Advancing_Neuromorphic_Computing_With_Loïhi_A_Survey_of_Results_and_Outlook

9. Arxiv今日论文 | 2025-11-21 http://lonepatient.top/2025/11/21/arxiv_papers_2025-11-21
10. Similarity Search and Applications - Springer Link <https://link.springer.com/content/pdf/10.1007/978-3-031-46994-7.pdf>
11. How do I validate JSON using an existing schema file in ... <https://stackoverflow.com/questions/44733603/how-do-i-validate-json-using-an-existing-schema-file-in-rust>
12. Autonomous and Intelligent Systems (AIS) Standards <https://standards.ieee.org/initiatives/autonomous-intelligence-systems/standards/>
13. Building an Open AIBOM Standard in the Wild <https://arxiv.org/html/2510.07070v1>
14. D7.8 - Report on the standardization landscape and ... <https://ec.europa.eu/research/participants/documents/downloadPublic?documentIds=080166e5db92e990&appId=PPGMS>
15. Artificial intelligence standardization roadmap <https://www.itu.int/epublications/publication/itu-t-y-suppl-72-2022-11-itu-t-y-3000-series-artificial-intelligence-standardization-roadmap>
16. IEEE Standard for Fail-Safe Design of Autonomous and ... https://www.researchgate.net/publication/394360658_IEEE_Standard_for_Fail-Safe_Design_of_Autonomous_and_Semi-Autonomous_Systems_Putting_Principles_Into_Practice
17. Hybrid GNN–LSTM defense with differential privacy and ... <https://www.nature.com/articles/s41598-025-27691-6>
18. Rust 在云原生中的应用：使用kube-rs 构建Kubernetes ... https://blog.csdn.net/m0_37613794/article/details/154146911
19. Rust 在云原生中的应用：使用`kube-rs 构建Kubernetes ... https://blog.csdn.net/LZL_SQL/article/details/154130893
20. Formal Verification for Preventing Misconfigured Access ... <https://ieeexplore.ieee.org/iel8/6287639/10820123/11122676.pdf>
21. RagVerus: Repository-Level Program Verification with ... <https://arxiv.org/html/2502.05344v1>
22. Anvil: verifying liveness of cluster management controllers <https://dl.acm.org/doi/10.5555/3691938.3691973>
23. Formal Verification for Preventing Misconfigured Access ... <https://ieeexplore.ieee.org/document/11122676/>
24. Admission Control in Kubernetes <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/>

25. (PDF) Z3: an efficient SMT solver https://www.researchgate.net/publication/225142568_Z3_an_efficient_SMT_solver
26. Don't Unwrap in Production: A Formal Verification Guide <https://dev.to/prasincs/dont-unwrap-in-production-a-formal-verification-guide-49ei>
27. Validating Admission Policy <https://kubernetes.io/docs/reference/access-authn-authz/validating-admission-policy/>
28. Kubernetes Validating Admission Policies: A Practical ... <https://kubernetes.io/blog/2023/03/30/kubescape-validating-admission-policy-library/>
29. Chapter 6. ValidatingAdmissionPolicyBinding ... https://docs.redhat.com/en/documentation/openshift_container_platform/4.17/html/extension_apis/validatingadmissionpolicybinding-admissionregistration-k8s-io-v1
30. Evolution of the IEEE P7009 Standard: Towards Fail-Safe ... https://www.researchgate.net/publication/363116347_Evolution_of_the_IEEE_P7009_Standard_Towards_Fail-Safe_Design_of_Autonomous_Systems
31. IEEE Standard for Fail-Safe Design of Autonomous and ... <https://zenodo.org/api/records/13292022/files/7009X-EBI.pdf/content>
32. Artificial Intelligence (RP2024) | Interoperable Europe Portal <https://interoperable-europe.ec.europa.eu/collection/rolling-plan-ict-standardisation/artificial-intelligence-rp2024>
33. STA N D A R D S https://img.antpedia.com/standard/files/pdfs_ora/20240709/IEEE%20Std%207009-2024.pdf
34. IEEE 7010: A New Standard for Assessing the Well-being ... <https://arxiv.org/pdf/2005.06620>
35. Extend the Kubernetes API with CustomResourceDefinitions <https://kubernetes.io/docs/tasks/extend-kubernetes/custom-resources/custom-resource-definitions/>
36. Model Context Protocol servers <https://gitee.com/mirrors/Model-Context-Protocol>
37. Towards Repository-Level Program Verification with Large ... <https://www.arxiv.org/pdf/2509.25197>
38. Chapter 5. ValidatingAdmissionPolicy ... https://docs.redhat.com/en/documentation/openshift_container_platform/4.17/html/extension_apis/validatingadmissionpolicy-admissionregistration-k8s-io-v1
39. 平台和工具 - 架构师研究会 <https://architect.pub/book/export/html/36>
40. Brew Formula | PDF | Command Line Interface <https://www.scribd.com/document/551715438/Brew-Formula>
41. Enforce Pod Security Standards by Configuring the Built-in ... <https://kubernetes.io/docs/tasks/configure-pod-container/enforce-standards-admission-controller/>

42. Dynamic Admission Control <https://kubernetes.io/docs/reference/access-authn-authz/extensible-admission-controllers/>
43. Study About the Capes Portal of E-Journals Non-users https://www.academia.edu/26931011/Studу_About_the_Capes_Portal_of_E_Journals_Non_users
44. How to create JSON object strategy according to a schema ... <https://stackoverflow.com/questions/72986320/how-to-create-json-object-strategy-according-to-a-schema-with-rust-proptest>
45. Icsea 2022 Full | PDF | Data <https://www.scribd.com/document/969442331/Icsea-2022-Full>
46. Neural Coding in Spiking Neural Networks: A Comparative ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC7970006/>
47. Evaluating Spiking Neural Network On Neuromorphic ... <https://arxiv.org/pdf/2308.00787>
48. Spiking Neural Networks for Multimodal Neuroimaging <https://pmc.ncbi.nlm.nih.gov/articles/PMC12189790/>
49. Exploring Neuromorphic Computing Based on Spiking ... https://www.researchgate.net/publication/365479857_Exploring_Neuromorphic_Computing_Based_on_Spiking_Neural_Networks_Algorithms_to_Hardware
50. Neuromorphic Spiking Neural Networks and Their ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC6747825/>
51. Unleashing the Power of Kubernetes 1.26 <https://dev.to/douglasmakey/cel-for-admission-controller-with-validatingadmissionpolicy-in-k8s-126-520b>
52. Syntaxe de flux de travail pour GitHub Actions <https://docs.github.com/fr/actions/reference/workflows-and-actions/workflow-syntax>
53. Communications202108 DL | PDF <https://www.scribd.com/document/595629726/communications202108-dl>
54. A Survey of Agent Interoperability Protocols: Model Context ... <https://arxiv.org/html/2505.02279v2>
55. 小海/awesome-mcp-servers <https://gitee.com/boomer001/awesome-mcp-servers/blob/main/README.md>
56. Пакети в „sid“, Под-раздел python <https://packages.debian.org/bg/sid/python/>
57. Spiking Neural Network Integrated Circuits <https://arxiv.org/pdf/2203.07006>
58. Ethics of neurotechnology: UNESCO adopts the first global ... <https://www.unesco.org/en/articles/ethics-neurotechnology-unesco-adopts-first-global-standard-cutting-edge-technology>

59. The Ethics of Neurotechnology: UNESCO appoints ... <https://www.unesco.org/en/articles/ethics-neurotechnology-unesco-appoints-international-expert-group-prepare-new-global-standard>
60. Report of the International Bioethics Committee ... <https://unesdoc.unesco.org/ark:/48223/pf0000378724>
61. RFC 8610: 3 of 4, p. 36 to 52 <https://www.tech-invite.com/y85/tinv-ietf-rfc-8610-3.html>
62. Blog <https://www.akamai.com/blog?ref=ecommerceberlin&page=66>
63. Using Z3 for Formal Modeling and Verification of FNN ... <https://arxiv.org/abs/2304.10558>
64. Kubernetes admission control with validating webhooks <https://developers.redhat.com/articles/2021/09/17/kubernetes-admission-control-validating-webhooks>
65. UNESCO Adopts First Global Framework on ... <https://www.globalpolicywatch.com/2026/01/unesco-adopts-first-global-framework-on-neurotechnology-ethics/>
66. Admission Webhook Good Practices <https://kubernetes.io/docs/concepts/cluster-administration/admission-webhooks-good-practices/>
67. The 2 main parts of a Kubernetes Validating Admission ... <https://dev.to/muaazsaleem/the-2-main-parts-of-a-kubernetes-validating-admission-webhook-11ga>
68. Draft Recommendation on the Ethics of Neurotechnology <https://unesdoc.unesco.org/ark:/48223/pf0000394866>
69. Final report on the draft Recommendation on the Ethics of ... <https://unesdoc.unesco.org/ark:/48223/pf0000393266>
70. First draft of the Recommendation on the Ethics ... <https://unesdoc.unesco.org/ark:/48223/pf0000391444>
71. Draft text of the Recommendation on the Ethics ... <https://unesdoc.unesco.org/ark:/48223/pf0000393395>
72. First draft of a Recommendation on the Ethics ... <https://unesdoc.unesco.org/ark:/48223/pf0000391074>
73. (PDF) The UNESCO draft Recommendations on ethics of ... https://www.researchgate.net/publication/391196602_The_UNESCO_draft_Recommendations_on_ethics_of_Neurotechnology_-A_commentary
74. Using Z3 for Formal Modeling and Verification of FNN ... <https://arxiv.org/pdf/2304.10558.pdf>
75. unesco <https://research.bjmu.edu.cn/docs/2024-10/af391e2749cf4794b652945bebbb5b91.pdf>

76. Final report on the draft text of the Recommendation ... <https://unesdoc.unesco.org/ark:/48223/pf0000393400>
77. Recommendation on the ethics of neurotechnology <https://www.unesco.org/en/node/86248>
78. Ethics of neurotechnology <https://www.unesco.org/en/ethics-neurotech>
79. Intergovernmental Meeting on the draft Recommendation ... <https://www.unesco.org/en/articles/intergovernmental-meeting-draft-recommendation-ethics-neurotechnology>
80. Towards Efficient Formal Verification of Spiking Neural ... <https://arxiv.org/html/2408.10900v1>
81. Modeling and Verification of Sigma Delta Neural Networks ... <https://dl.acm.org/doi/pdf/10.1145/3735452.3735525>
82. A Sampling-based Method for Output Safety Verification ... https://www.researchgate.net/publication/395087297_A_Sampling-based_Method_for_Output_Safety_Verification_of_Spiking_Neural_Networks
83. Verifying Quantized Neural Networks using SMT-Based ... <https://arxiv.org/abs/2106.05997>
84. Configuring Safe Spiking Neural Controllers for Cyber- ... https://www.researchgate.net/publication/382884845_Configuring_Safe_Spiking_Neural.Controllers_for_Cyber-Physical_Systems_through_Formal_Verification
85. Linear leaky-integrate-and-fire neuron model based spiking ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC9448910/>
86. Mapping Functional Connectivity from the Dorsal Cortex to ... <https://www.sciencedirect.com/science/article/pii/S089662732030492X>
87. Securing Admission Controllers <https://kubernetes.io/blog/2022/01/19/secure-your-admission-controllers-and-webhooks/>
88. 11 Kubernetes Admission Controller Best Practices for ... <https://www.redhat.com/en/blog/11-kubernetes-admission-controller-best-practices-for-security>
89. A Guide to Kubernetes Admission Controllers <https://kubernetes.io/blog/2019/03/21/a-guide-to-kubernetes-admission-controllers/>
90. Analysis of Neuronal Spike Trains, Deconstructed - PMC - NIH <https://pmc.ncbi.nlm.nih.gov/articles/PMC4970242/>
91. Verification of a neuromorphic computing network ... <https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2022.958343/full>
92. Advancements in neuromorphic computing for bio-inspired ... <https://www.sciencedirect.com/science/article/pii/S0925231225018934>

93. (PDF) Proof-of-Spiking-Neurons(PoSN): Neuromorphic ... https://www.researchgate.net/publication/397322038_Proof-of-Spiking-NeuronsPoSN_Neuromorphic_Consensus_for_Next-Generation_Blockchains
94. Algorithm-architecture adequacy of spiking neural ... <https://theses.hal.science/tel-02400657v1/file/2018TOU30318a.pdf>
95. From Brains to Systems <https://link.springer.com/content/pdf/10.1007/978-1-4614-0164-3.pdf>
96. (PDF) First-spike based visual categorization using reward- ... https://www.researchgate.net/publication/317164116_First-spike_based_visual_categorization_using_reward-modulated_STDP
97. A Multi-Language Stack for Emotional Avatar Generation https://www.linkedin.com/posts/william-allen-3381a7312_2-smart-glasses-python-streams-biometric-activity-7386400661296521216-2CYd
98. Latest articles https://www.mdpi.com/latest_articles
99. Invariant neural dynamics drive commands to control ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC10527529/>
100. Cycle is All You Need: More Is Different <https://www.arxiv.org/pdf/2509.21340>
101. Diverse and flexible behavioral strategies arise in recurrent ... <https://journals.plos.org/ploscompbiol/article?id=10.1371/journal.pcbi.1013559>
102. The Case for Kubernetes Resource Limits: Predictability vs. ... <https://kubernetes.io/blog/2023/11/16/the-case-for-kubernetes-resource-limits/>
103. Project 3 Data Science Thought Leadership (Final) https://www.rpubs.com/Rajwantmishra/project_spark_607_3final
104. Telefonica Tech · Blog <https://telefonicatech.com/en/blog/author/telefonicatech>
105. Spiking representation learning for associative memories <https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2024.1439414/full>
106. Repetitive training enhances the pattern recognition ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12064202/>
107. A visual encoding model links magnetoencephalography ... <https://www.sciencedirect.com/science/article/pii/S1053811921009289>
108. Long-range functional loops in the mouse olfactory system ... [https://www.cell.com/neuron/fulltext/S0896-6273\(22\)00810-8](https://www.cell.com/neuron/fulltext/S0896-6273(22)00810-8)
109. A Massively Parallel CRISPR-Based Screening Platform for ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC11844385/>
110. Roadmap on biology in time varying environments - IOPscience <https://iopscience.iop.org/article/10.1088/1478-3975/abde8d>

111. Canonical circuit computations for computer vision <https://link.springer.com/article/10.1007/s00422-023-00966-9>
112. Invariant neural dynamics drive commands to control ... <https://www.sciencedirect.com/science/article/pii/S0960982223007789>
113. Invariant neural dynamics drive commands to control ... [https://www.cell.com/current-biology/pdf/S0960-9822\(23\)00778-9.pdf](https://www.cell.com/current-biology/pdf/S0960-9822(23)00778-9.pdf)
114. What should be on a checklist that would help someone ... <https://stackoverflow.com/questions/915649/what-should-be-on-a-checklist-that-would-help-someone-develop-good-oo-software>
115. Visual prototypes in the ventral stream are attuned to ... <https://www.nature.com/articles/s41467-021-27027-8>
116. A machine learning framework to optimize optic nerve ... <https://www.sciencedirect.com/science/article/pii/S2666389921001197>
117. Retinal processing: insights from mathematical modelling <https://inria.hal.science/hal-03454859v2/document>
118. Retinal processing: insights from mathematical modelling <https://arxiv.org/pdf/2201.06795>
119. Advancements in neural network acceleration <https://www.sciencedirect.com/science/article/pii/S2405959525001687>
120. AI-Driven Autonomous Cyber-Security Systems: Advanced ... <https://www.linkedin.com/pulse/ai-driven-autonomous-cyber-security-systems-advanced-ramachandran-lmame>
121. Multi-Partner project: dAIEDGE - A network of excellence ... https://hal.science/hal-05464162v1/file/8012_pdf_upload.pdf
122. Network Digital Twin for 6G and Beyond: An End-to- ... <https://ieeexplore.ieee.org/iel8/8782661/10829557/11130513.pdf>
123. c# - Validation strategies <https://stackoverflow.com/questions/6599501/validation-strategies>
124. Retinal processing: insights from mathematical modelling <https://inria.hal.science/hal-03454859/document>
125. Chemical Imaging for Biological Systems: Techniques, AI ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12208185/>
126. ROSE: A Neurocomputational Architecture for Syntax - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC10055479/>
127. java - Why can I use an interface as a type when creating ... <https://stackoverflow.com/questions/24742024/why-can-i-use-an-interface-as-a-type-when-creating-a-reference-to-an-object>

128. Computer Science <https://arxiv.org/list/cs/new>
129. A Review of Abstraction Methods Toward Verifying Neural ... <https://dl.acm.org/doi/full/10.1145/3617508>
130. A Survey on Formal Verification Techniques for Safety- ... <https://www.mdpi.com/2079-9292/7/6/81>
131. A Survey of Neural Network Robustness Assessment in ... <https://arxiv.org/pdf/2404.08285>
132. UCLA <https://escholarship.org/content/qt9wv5s2j8/qt9wv5s2j8.pdf>